

Model output report: Isolation_Forest_c05_n100

Summary of model_output_Isolation_Forest_c05_n100.csv (fusion of rule-based + Isolation_Forest_c05_n100 scores).

1. Overview

Item	Value
Total customers	61,410
Flagged (top 5%)	3,071 (5.00%)
Not flagged	58,339
Effective threshold (risk_score $\geq$)	0.3934

Insight: This model flags the top 5% of customers by combined risk score (rule-based plus Isolation_Forest_c05_n100 anomaly score), so 3,071 customers are marked for review. The threshold 0.3934 is the minimum score needed to be in that top 5%. All flagged customers should be prioritised for further investigation; the rest of the population is below the cutoff.

2. Risk score distribution

Statistic	All	Flagged	Not flagged
Min	0.0230	0.3935	0.0230
Max	0.7642	0.7642	0.3934
Mean	0.2342	0.4515	0.2228

Insight: The spread of scores (min to max, and standard deviation) shows how much the model separates customers. Flagged customers have higher mean and max scores than the rest; the gap between flagged and not-flagged means indicates how distinct the high-risk group is. IQR (interquartile range) is 0.110, so the middle 50% of scores span this range; a wider IQR suggests more variation in risk across the population.

3. Distribution by score band

Score band	Count	% of total
[0.0, 0.1)	1,954	3.2%
[0.1, 0.2)	21,324	34.7%
[0.2, 0.3)	25,860	42.1%
[0.3, 0.4)	9,528	15.5%
[0.4, 0.5)	2,249	3.7%
[0.5, 0.6)	415	0.7%
[0.6, 0.7)	74	0.1%

[0.7, 0.8)	6	0.0%
[0.8, 0.9)	0	0.0%
[0.9, 1.0)	0	0.0%

Insight: Most customers fall in the lower score bands; only a small share have scores in the 0.5+ range (0.8% here). The top 5% flagged set sits at the right tail of this distribution. If there are customers just below the threshold (within 0.02), there are 1,192 in that borderline band—these may warrant monitoring as small changes in behaviour could push them into the review set.

4. How to use this output

Use the corresponding `model_output_*.csv` and `model_output_*_explanations.csv` for Task 2 and Task 3. The explanations file gives a plain-English reason for each customer's score; the report above summarises the overall distribution and helps you understand how this model's fusion (rule + anomaly) behaves across the population.